Lab Instructions - session 5

Reading from camera device, edge detection

Part 1. Reading from a webcam

The following code reads a video stream from a webcam and displays it.
File: **read_webcam.py**

```python
import numpy as np
import cv2

cam_id = 0  # camera id

# for default webcam, cam_id is usually 0
# try out other numbers (1,2,..) if this does not work

cap = cv2.VideoCapture(cam_id)

while True:
    ret, I = cap.read();

    cv2.imshow("my stream", I);

    # press "q" to quit
    if cv2.waitKey(1) & 0xFF == ord('q'):
        break

cap.release()
cv2.destroyAllWindows()
```

Part 2. Edge detection

Computing the Gradients

We compute Sobel gradients both by applying a Sobel filter and by using the OpenCV function "Sobel".

File: **sobel1.py**

```python
import numpy as np
import cv2
from matplotlib import pyplot as plt

I = cv2.imread("agha-bozorg.jpg", cv2.IMREAD_GRAYSCALE)

# Compute the gradient in x direction using the sobel filter

# Method 1: using filter2D **********
Dx = np.array([[-1, 0, 1],
               [-2, 0, 2],
               [-1, 0, 1]]); # Sobel filter

Ix = cv2.filter2D(I, -1, Dx);
print(I.dtype)
print(Ix.dtype)

Ix = cv2.filter2D(I, cv2.CV_16S, Dx); # cv2.CV_16S: 16 bit signed integer
print(Ix.dtype)
input('press ENTER to continue... ')

# Method 2: using sobel function **********
Ix2 = cv2.Sobel(I,cv2.CV_16S,1,0)
print(np.abs(Ix - Ix2).max())
input('press ENTER to continue... ')

# Plot the gradient image
f, axes = plt.subplots(2, 2)

axes[0,0].imshow(I,cmap = 'gray')
axes[0,0].set_title("Original Image")

axes[0,1].imshow(Ix,cmap = 'gray')
axes[0,1].set_title("Ix (cv2.filter2D)")

axes[1,0].imshow(Ix2,cmap = 'gray')
axes[1,0].set_title("Ix2 (cv2.Sobel)")

axes[1,1].imshow(np.abs(Ix),cmap = 'gray')
axes[1,1].set_title("abs(Ix)")

# Notice that imshow in matplotlib considers the minimums value of I
# as black and the maximum value as white (this is different from
# the behavior in cv2.imshow
plt.show()
```

- Why have we used `cv2.CV_16s` (16 bit signed integer) format for the output of filter2D and Sobel functions, instead of -1 (which gives the same numeric type as the input image, i.e. CV_8U or unsigned 8-bit integer)?
- What is the value of `np.abs(Ix - Ix2).max()`? Are `Ix` and `Ix2` different?
- Why are most of the pixels in images of `axes[0,1]` and `axes[1,0]` gray? What do the black and white pixels show in these images? (notice that matplotlib automatically sets the minimum value of an image to black and the maximum value to white)

Towards edge detection!

The strength of the edge at each certain pixel can be represented by the magnitude of the gradient vector, that is

$$\sqrt{(\partial I/\partial x)^2 + (\partial I/\partial y)^2}$$

The next code plots the gradient in x and y directions, plus the magnitude of the gradient.

File: **sobel2.py**

```python
import numpy as np
import cv2
from matplotlib import pyplot as plt

I = cv2.imread("agha-bozorg.jpg", cv2.IMREAD_GRAYSCALE)

# Sobel gradient in the x direction
Ix = cv2.Sobel(I,cv2.CV_64F,1,0)
print(Ix.dtype)

# Sobel gradient in the y direction
Iy = cv2.Sobel(I,cv2.CV_64F,0,1)
print(Iy.dtype)

# Magnitude of gradient
E = np.sqrt(Ix*Ix + Iy*Iy)

# Plot the gradient image
f, axes = plt.subplots(2, 2)

axes[0,0].imshow(I,cmap = 'gray')
axes[0,0].set_title("Original Image")

axes[0,1].imshow(abs(Ix),cmap = 'gray')
axes[0,1].set_title("abs(Ix)")


axes[1,0].imshow(abs(Iy),cmap = 'gray')
axes[1,0].set_title("abs(Iy)")

axes[1,1].imshow(E,cmap = 'gray')
axes[1,1].set_title("Magnitude of Gradient")
plt.show()
```

Edge detection (Sobel, Laplacian, Canny):

Now, we move on to edge detection. We use three different methods:
1. Using Sobel gradients + thresholding
2. Laplacian of Gaussian (LoG) + detection zero-crossings
3. Canny Edge detection

The function **zero_crossing** finds zero crossings in an image for LoG edge detection. You do not need to know how the functions **std_filter** and **zero_crossing** work.

File: **edge.py**

```python
import numpy as np
import cv2
from matplotlib import pyplot as plt

def std_filter(I, ksize):
    F = np.ones((ksize,ksize), dtype=np.float) / (ksize*ksize);

    MI = cv2.filter2D(I,-1,F) # apply mean filter on I

    I2 = I * I; # I squared
    MI2 = cv2.filter2D(I2,-1,F) # apply mean filter on I2

    return np.sqrt(MI2 - MI * MI)

def zero_crossing(I):
    """Finds locations at which zero-crossing occurs, used for
    Laplacian edge detector"""

    Ishrx = I.copy();
    Ishrx[:,1:] = Ishrx[:,:-1]

    Ishdy = I.copy();
    Ishdy[1:,:] = Ishdy[:-1,:]

    ZC = (I==0) | (I * Ishrx < 0) | (I * Ishdy < 0); # zero crossing
locations

    SI = std_filter(I, 3) / I.max()

    Mask =  ZC & (SI > .1)

    E = Mask.astype(np.uint8) * 255 # the edges

    return E

I = cv2.imread("agha-bozorg.jpg", cv2.IMREAD_GRAYSCALE)

# set the sigma for Gaussian Blurring
sigma = 7

# Sobel magnitude of gradient
thresh = 90 # threshold
```

```python
Ib = cv2.GaussianBlur(I, (sigma,sigma), 0); # blur the image
Ix = cv2.Sobel(Ib,cv2.CV_64F,1,0)
Iy = cv2.Sobel(Ib,cv2.CV_64F,0,1)
Es = np.sqrt(Ix*Ix + Iy*Iy)
Es = np.uint8(Es > thresh)*255 # threshold the gradients

# Laplacian of Gaussian
# Here, we first apply a Gaussian filter and then apply
# the Laplacian operator (instead of applying the LoG filter)
Ib = cv2.GaussianBlur(I, (sigma,sigma), 0);
El = cv2.Laplacian(Ib,cv2.CV_64F,ksize=5)
El = zero_crossing(El);

# Canny Edge detector
lth = 50    # low threshold
hth = 120   # high threshold
Ib = cv2.GaussianBlur(I, (sigma,sigma), 0); # blur the image
Ec = cv2.Canny(Ib,lth, hth)
f, axes = plt.subplots(2, 2)

axes[0,0].imshow(I,cmap = 'gray')
axes[0,0].set_title("Original Image")

axes[0,1].imshow(Es,cmap = 'gray')
axes[0,1].set_title("Sobel")


axes[1,0].imshow(El,cmap = 'gray')
axes[1,0].set_title("Laplacian")

axes[1,1].imshow(Ec,cmap = 'gray')
axes[1,1].set_title("Canny")

# Notice that imshow in matplotlib considers the minimums value of I
# as black and the maximum value as white (this is different from
# the behavior in cv2.imshow)
plt.show()
```

- Compare the Canny edge detector to Sobel+thresholding. Can you see the effect of non-maximum suppression?
- Notice that in all cases we first smooth the image using a Gaussian filter. What is the purpose of smoothing the image? Change the smoothing parameter **sigma** (from **sigma=7** to sigma=5, 3, 1 or 9) and see what happens by increasing or decreasing this parameter.
- Change the low and high thresholds of the Canny edge detector and observe the results.

Today's task:

You need to read a video stream from your webcam and apply different gradients or edge detection operations to the stream. You have to do this by completing the file **lab5_task1.py.** Your program must have the following functionalities:

- **press the 'o' key:** show the original webcam frame (already done)
- **press the 'x' key:** show the Sobel gradient in the x direction (already done)
- **press the 'y' key:** show the Sobel gradient in the y direction
- **press the 'm' key:** show the Sobel magnitude of the gradient
- **press the 's' key:** show the result of Sobel + thresholding edge detection
- **press the 'l' key:** apply Laplacian of Gaussian (LoG) edge detector
- **press the 'c' key:** apply Canny edge detector
- **press the '+' key:** increase smoothing parameter sigma (already done)
- **press the '-' key:** decrease smoothing parameter sigma (already done)
- **press the 'q' key:** quit the program (already done)

Notice that after reading each image frame we apply a Gaussian filter to blur it. All gradient/edge detection operations must be done on the blurred image `Ib`.

File: **lab5_task1.py**

```python
import numpy as np
import cv2


cam_id = 0  # camera id

# for default webcam, cam_id is usually 0
# try out other numbers (1,2,..) if this does not work

cap = cv2.VideoCapture(cam_id)

mode = 'o' # show the original image at the beginning

sigma = 5

while True:
    ret, I = cap.read();
    #I = cv2.imread("agha-bozorg.jpg") # can use this for testing
    I = cv2.cvtColor(I, cv2.COLOR_BGR2GRAY) # convert to grayscale
    Ib = cv2.GaussianBlur(I, (sigma,sigma), 0); # blur the image

    if mode == 'o':
        # J = the original image
        J = I
    elif mode == 'x':
        # J = Sobel gradient in x direction
        J = np.abs(cv2.Sobel(Ib,cv2.CV_64F,1,0));
```

```python
    elif mode == 'y':
        # J = Sobel gradient in y direction
        pass
    elif mode == 'm':
        # J = magnitude of Sobel gradient
        pass
    elif mode == 's':
        # J = Sobel + thresholding edge detection
        pass
    elif mode == 'l':
        # J = Laplacian edges
        pass
    elif mode == 'c':
        # J = Canny edges
        pass

    # we set the image type to float and the maximum value to 1
    # (for a better illustration) notice that imshow in opencv does not
    # automatically map the min and max values to black and white.
    J = J.astype(np.float) / J.max();
    cv2.imshow("my stream", J);

    key = chr(cv2.waitKey(1) & 0xFF)

    if key in ['o', 'x', 'y', 'm', 's', 'c', 'l']:
        mode = key
    if key == '-' and sigma > 1:
        sigma -= 2
        print("sigma = %d"%sigma)
    if key in ['+','=']:
        sigma += 2
        print("sigma = %d"%sigma)
    elif key == 'q':
        break

cap.release()
cv2.destroyAllWindows()
```

- Wave your hand quickly. What happens to the edges? Why?
- Press the 's' and 'c' keys to compare **Canny** with **Sobel** edge detection.
  Explain the effect of non-maximum suppression to the TA.
- In each case, increase and decrease the **sigma** by pressing + and - keys and
  see what happens. Explain the reason to the TA.

References

- https://docs.opencv.org/3.0-beta/doc/py_tutorials/py_imgproc/py_gradients/py_gradients.html#gradients
- https://docs.opencv.org/3.0-beta/doc/py_tutorials/py_imgproc/py_canny/py_canny.html#canny